Cekgu

The answer key is wrong. Nobody checked.

Cekgu reviews a quiz before learners sit it.

- Two models, read blind
- Every reading receipted
- Refuses on disagreement

TEAM M1KU

@AlaskanTuna Gateway · Queue · API	@kymil4 Frontend · Deploy	@chaosiris Testing · Deck	@c3638 Docs · Submission

1. PROBLEM STATEMENT

One wrong key. Every learner pays.

The damaging defects are small, and every check that exists today either anchors on the author’s own key or runs after the marks are out.

THE EDUCATOR, TODAY

Publishes a multiple-choice paper several times a month. Subject-competent, time-poor, and **personally accountable when an item is challenged.**

Writes	Twelve items, each with a keyed answer.
Self-reviews	Fast, but anchored to the answer they intended.
Asks a colleague	Useful, and it spends an expert’s time on every item.
Publishes	Learners sit it.
Gets a challenge	One learner disputes an item. Investigate, remark, communicate, withdraw.

A composite of the primary operator in PRODUCT.md, not a named person.

WHAT A DEFECT LOOKS LIKE

Which data structure removes elements in first in, first out order?

SUPPLIED KEY	BOTH BLIND READERS	
A · Stack	B · Queue	Possible Key Error

WHAT EXISTS TODAY, AND THE GAP IT LEAVES

Self-review	Anchored to the author’s intended answer.
A colleague re-solves it	Scarce subject-expert time, spent on every item.
One general AI chat	One opaque opinion, may anchor on the key, no proof of two independent readings.
Enterprise platform	Authoring and workflow suites, heavy for one educator who needs a preflight.
Post-exam item analysis	Repairs marks later. Cannot prevent the first unfair result.

Human moderation already exists because the risk is real. **UiTM’s final-examination guidance requires a faculty vetting committee.** Cekgu is the pass before that committee, not a replacement for it.

2. PROJECT OBJECTIVE

Spend human attention where it matters.

Cekgu re-solves every item so the educator reads only the ones that came back different.

CEKGU RE-SOLVES EVERY ITEM, BLIND **12**

1 What is the time complexity of reading the element at a known ind...
2 Which transport protocol retransmits lost segments so that data a...
3 Which data structure removes elements in first in, first out orde...
4 In big-O notation, what does O(n) say about an algorithm?
5 Mutual exclusion, hold and wait, and no preemption are three of t...
6 Which layer of the TCP/IP model does HTTP belong to?
7 What does a primary key guarantee about a relational table?
8 What does a compiler do with a source program?
9 What is the main job of the Domain Name System?
10 Why does a processor have a cache?
11 How many bytes are there in one kilobyte?
12 What does creating a branch in Git let a developer do?

SORTED →

ONLY THESE NEED A HUMAN **3**

3 Which data structure removes elements in firs... — Possible Key Error

9 What is the main job of the Domain Name Syste... — Possible Key Error

11 How many bytes are there in one kilobyte? — Possible Ambiguity

9 items — Clear

PUBLIC SAMPLE RECORD · 12 ITEMS

3. OVERALL CONCEPT

Two readers, blind. Then a fixed rule.

Neither model sees the key or the other’s answer. A fixed rule compares the two readings first, and the key only after.

THE ITEM AS SUBMITTED

Which data structure removes elements in first in, first out order?

Your key is withheld.

→

TWO MODEL FAMILIES, ANSWERING BLIND

MiniMax-M2.7 — Queue

=

Kimi-K2.6 — Queue

≠

THE SUPPLIED KEY — Stack

→

CHECKED IN THIS ORDER

1. **Unverified**
   fewer than two verified readings
2. **Split Opinion**
   the two answers differ
3. **Possible Ambiguity**
   both list more than one option
4. **Clear**
   the shared answer is the key
5. **Possible Key Error**
   the shared answer is not the key

3. OVERALL CONCEPT · THE MOMENT

Nobody told them the key. Both chose Queue.

Question 3 of the public sample record, exactly as it came back.

Which data structure removes elements in first in, first out order?

SUPPLIED KEY

A B C D

Stack

MiniMaxAI/MiniMax-M2.7

A B C D

Queue

RECEIPT Verified

moonshotai/Kimi-K2.6

A B C D

Queue

RECEIPT Verified

Possible Key Error

Both readers chose Queue. The supplied key is Stack.

The supplied key says Stack.

Cekgu asks a human to look. It changes nothing.

RULE FIRED

Two verified readings agree on a non-key option.

3. OVERALL CONCEPT · THE PROOF

Every reading carries a request id.

Each reading is tied to a GonkaRouter receipt anyone can fetch. A reading without one never enters a verdict.

PUBLIC SAMPLE RECORD · QUESTION 9

SERVED MODEL

MiniMaxAI/MiniMax-M2.7

REQUEST ID

req-1788427238422211326-414866

RECEIPT

Verified

SERVED MODEL

moonshotai/Kimi-K2.6

REQUEST ID

req-1788427238422211356-414867

RECEIPT

Verified

ASK THE GATEWAY YOURSELF

api.gonkarouter.io/v1/receipts/
req-1788427238422211326-414866

No key, no login. Anyone in this room can open it now.

Both ids are on the public sample record. Neither reading counts without one.


```
"model": "MiniMaxAI/MiniMax-M2.7",
"outcome": "success",
"status_code": 200,
"duration_ms": 30867,
"created_at": "2026-09-03T09:21:09Z",
"x_devshard_id": "70340"
```

4. TECHNOLOGY AND TRACK

One process. One gateway. Two models.

One Bun process, one gateway, four calls in flight, and a rule you can read off the screen.

1
Submit
Validated locally. The record exists in under a second, before any model runs.

2
Queue
Postgres. One item claimed at a time with `SKIP LOCKED`. Four gateway calls in flight.

3
Two blind readers
Two model families through GonkaRouter. Key withheld. `X-Gonka-No-Fallback` on every call.

4
Receipt check
`GET /v1/receipts/{id}`. The served model must equal the one asked for, within 5 s.

5
Fixed rule
Five verdicts, checked in order. Fewer than two distinct verified readings is Unverified.

6
Record
Verdict, every attempt, every request id. The human decision is stored beside it, never over it.

STACK

Browser	React 19, Vite
Server	Bun, Hono, worker in the same process
Data	Postgres on Neon, Drizzle
Host	Cloud Run

WHEN THE GATEWAY IS SLOW

Hedge at 45 s	a duplicate call to the same model, first answer wins
Cutoff at 90 s	the call is aborted and recorded with no request id
3 attempts per family	then the item is Unverified, never guessed

ASSERTED EVERY RUN

`only-gonkarouter.test.ts`

No other AI provider exists anywhere in the code. A judge can grep for it.

✓ All inference on GonkaRouter ✓ Two models, distinct receipts ✓ A request id per step ✓ The rule printed on screen

5. MOTIVATION AND CHALLENGES

No item had two readings in 30 seconds.

The gateway is slow and sometimes absent, so the design queues, hedges, retries and refuses rather than guesses.

So Cekgu queues the work, retries, and **fails closed.**

12 of 12

items, two verified readings each

PUBLIC SAMPLE RECORD

SECONDS TO A COMPLETED CALL · 3 SEP BENCHMARK

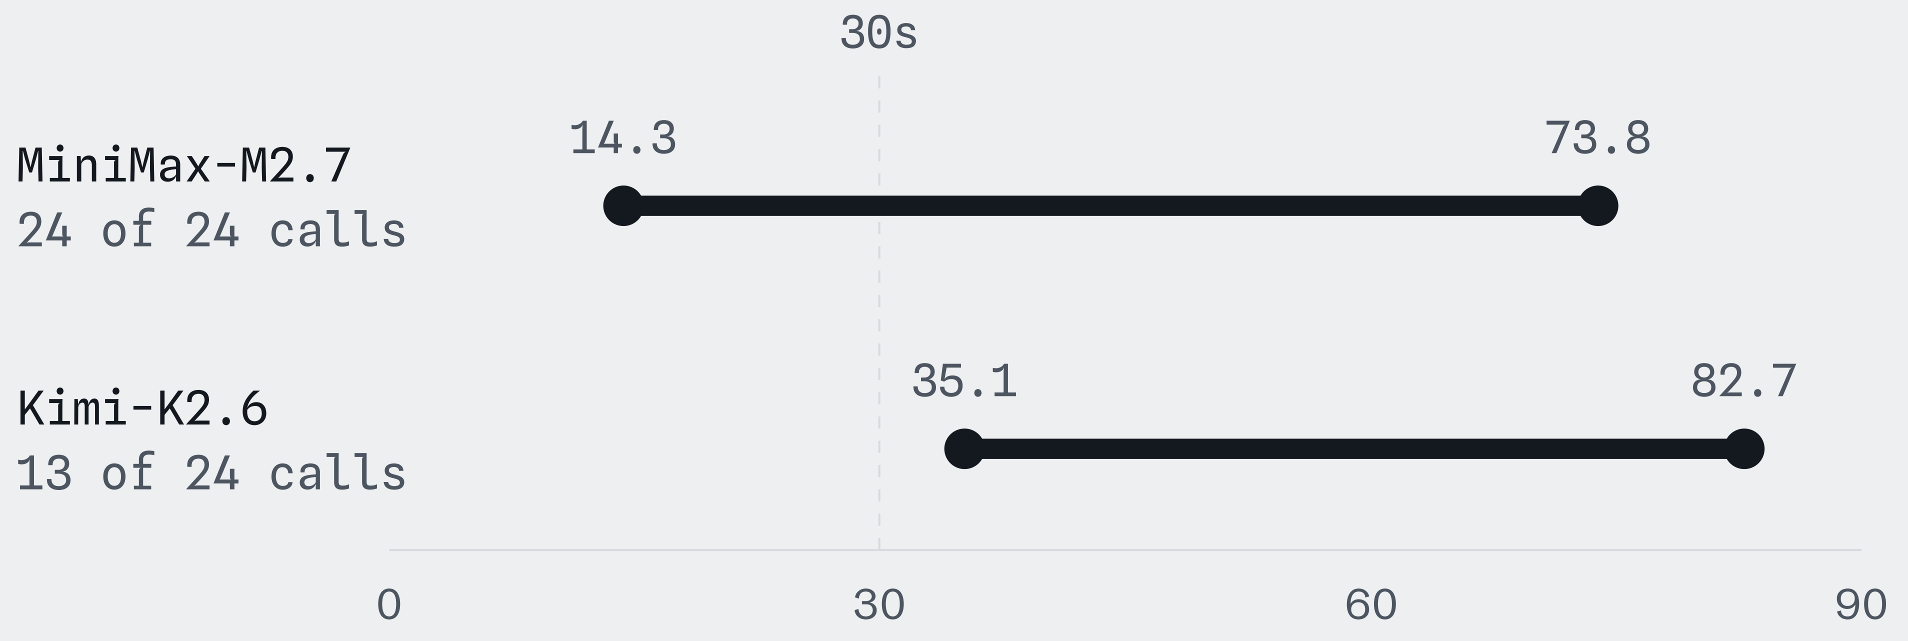

30-ITEM EVALUATION SET · RUN TWICE THROUGH THE DEPLOYED QUEUE

40 clean control runs

0 false flags

20 planted defect runs

14 caught

5 abstained as Unverified. 1 came back Clear.

6. BUSINESS MODEL

We charge for questions checked. Not tokens.

Priced by questions checked, so the educator never sees a model bill.

FREE	CEKGU PLUS	CEKGU STUDIO
20	300	1500
questions checked per month	questions checked per month	questions checked per month
RM0	RM29 / month	RM79 / month

The educator never sees a model bill. **Pilot pricing, and no buyer has paid yet.**

Guest is a shared demo workspace, expiring in 24 hours.

Check the exam before the exam checks us.

LIVE DEMO	cekgu-op7lf5dspq-as.a.run.app
REPOSITORY	github.com/MUBA-M1KU/Cekgu
DEVFOLIO	muba-hackathon.devfolio.co

Twelve items. Two blind readers each. Three sent back to a human.

@chaosiris · @AlaskanTuna · @kymil4 · @c3638

Scan to open the public sample record